

# Handling Data-skew Effects in Join Operations using MapReduce

M. Al Hajj Hassan<sup>1</sup>, M. Bamha<sup>2</sup>, and F. Loulergue<sup>2</sup>

<sup>1</sup> Lebanese International University, Beirut, Lebanon  
mohamad.hajjhassan01@liu.edu.lb

<sup>2</sup> Université Orléans, INSA Centre Val de Loire, LIFO EA 4022, France  
{mostafa.bamha,frederic.loulergue}@univ-orleans.fr

## Abstract

For over a decade, MapReduce has become a prominent programming model to handle vast amounts of raw data in large scale systems. This model ensures scalability, reliability and availability aspects with reasonable query processing time. However these large scale systems still face some challenges: data skew, task imbalance, high disk I/O and redistribution costs can have disastrous effects on performance.

In this paper, we introduce *MRFA-Join* algorithm: a new frequency adaptive algorithm based on MapReduce programming model and a randomised key redistribution approach for join processing of large-scale datasets. A cost analysis of this algorithm shows that our approach is insensitive to data skew and ensures perfect balancing properties during all stages of join computation. These performances have been confirmed by a series of experimentations.

*Keywords:* Join operations, Data skew, MapReduce model, Hadoop framework

## 1 Introduction

Today with the rapid development of network technologies, internet search engines, data mining applications and data intensive scientific computing applications, the need to manage and query a huge amount of datasets every day becomes essential. **Parallel processing of such queries on hundreds or thousands of nodes is obligatory to obtain a reasonable processing time [6].** However, building parallel programs on parallel and distributed systems is complicated because programmers must treat several issues such as load balancing and fault tolerance. **Hadoop [14] and Google's MapReduce model [8] are examples of such systems.** These systems are built from thousands of commodity machines and assure scalability, reliability and availability aspects [9]. **To reduce disk I/O, each file in such storage systems is divided into chunks or blocks of data and each block is replicated on several nodes for fault tolerance. Parallel programs are easily written on such systems following the MapReduce paradigm where a program is composed of a workflow of user defined *map* and *reduce* functions.**

Join operation is one of the most widely used operations in relational database systems, but it is also a heavily time consuming operation. For this reason it was a prime target for parallelization. The *join* of two relations  $R$  and  $S$  on attribute  $A$  of  $R$  and attribute  $B$  of  $S$  ( $A$  and  $B$  of the same domain) is the relation, written  $R \bowtie S$ , obtained by concatenating pairs of tuples from  $R$  and  $S$  for which  $R.A = S.B$ .

**Parallel join usually proceeds in two phases:** a redistribution phase (generally based on join attribute hashing and therefore called hashing algorithms) and then a sequential join of local fragments. Many parallel join algorithms have been proposed. The principal ones are: *Sort-merge join*, *Simple-hash join*, *Grace-hash join* and *Hybrid-hash join* [12]. All of them are based on hashing functions which redistribute relations such that all the tuples having the same join attribute value are forwarded to the same node. Local joins are then computed and their union is the output relation. Research has shown that join is parallelizable with near-linear speed-up on distributed architectures but only under ideal balancing conditions: data skew may have disastrous effects on the performance [13, 10]. To this end, several parallel algorithms were presented to handle data skew while treating join queries on parallel database systems [2, 3, 1, 13, 7, 10].

The aim of join operations is to combine information from two or more data sources. Unfortunately, MapReduce framework is somewhat inefficient to perform such operations since data from one source must be maintained in memory for comparison to other source of data. Consequently, adapting well-known join algorithms to MapReduce is not as straightforward as one might hope, and MapReduce programmers often use simple but inefficient algorithms to perform join operations especially in the presence of skewed data [11, 4, 9].

In [15], three well known algorithms for join evaluation were implemented using an extended MapReduce model. **These algorithms are *Sort-Merge-Join*, *Hash-Join* and *Block Nested-Loop Join*. Combining this model with a distributed file system facilitates the task of programmers because they don't need to take care of fault tolerance and load balancing issues.** However, load balancing in the case of join operations is not straightforward in the presence of data-skew. In [4] Blanas et al. have presented an improved versions of MapReduce sort-merge joins and semi-join algorithms for log processing to fix the problem of buffering all records from both inner and outer relations. For the same reasons as in parallel database management systems (PDBMS), even in the presence of integrated functionality for load balancing and fault tolerance in MapReduce, these algorithms still suffer from the effect of data skew. Indeed all the tuples having the same values in map phase are sent to the same reducer which limits the scalability of the presented algorithms [9].

In this paper we are interested in the evaluation of join operations on large scale systems using MapReduce. To avoid the effect of data skew, we introduce the MapReduce Frequency Adaptive Join algorithm (*MRFA-Join*) based on distributed histograms and a randomised key redistribution approach. This algorithm, inspired from our previous research on join and semi-join operations in PDBMS, is well adapted to manage huge amount of data on large scale systems even for highly skewed data. The remaining of the paper is organised as follows. In section 2 we briefly present the MapReduce programming model. Section 3 is devoted to the *MRFA-Join* algorithm and its complexity analysis. Experiments presented in section 4 confirm the efficiency of our approach. We conclude and give further research directions in section 5.

## 2 The MapReduce Programming Model

MapReduce [6] is a simple yet powerful framework for implementing distributed applications without having extensive prior knowledge of issues related to data redistribution, task allocation

or fault tolerance in large scale distributed systems.

Google's MapReduce programming model presented in [6] is based on two functions: **map** and **reduce**, that the programmer is supposed to provide to the framework. These two functions should have the following signatures:

**map:**  $(k_1, v_1) \longrightarrow \text{list}(k_2, v_2),$   
**reduce:**  $(k_2, \text{list}(v_2)) \longrightarrow \text{list}(v_3).$

The user must write the **map** function that has two input parameters, a key  $k_1$  and an associated value  $v_1$ . Its output is a list of intermediate key/value pairs  $(k_2, v_2)$ . This list is partitioned by the MapReduce framework depending on the values of  $k_2$ , where all pairs having the same value of  $k_2$  belong to the same group.

The **reduce** function, that must also be written by the user, has two parameters as input: an intermediate key  $k_2$  and a list of intermediate values  $\text{list}(v_2)$  associated with  $k_2$ . It applies the user defined merge logic on  $\text{list}(v_2)$  and outputs a list of values  $\text{list}(v_3)$ .

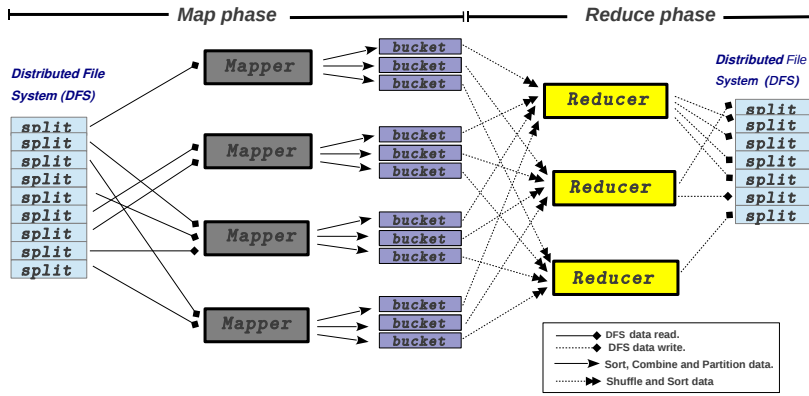


Figure 1: Map-reduce framework.

In this paper, we used an open source version of MapReduce called Hadoop developed by "The Apache Software Foundation". Hadoop framework includes a distributed file system called HDFS<sup>1</sup> designed to store very large files with streaming data access patterns.

For efficiency reasons, in Hadoop MapReduce framework, users may also specify a "Combine function", to reduce the amount of data transmitted from Mappers to Reducers during *shuffle* phase (see fig 1). The "Combine function" is like a local reduce applied (at map worker) before storing or sending intermediate results to the reducers. The signature of **combine** function is:

**combine:**  $(k_2, \text{list}(v_2)) \longrightarrow (k_2, \text{list}(v_3)).$

To cover a large range of applications needs in term of computation and data redistribution, in Hadoop framework, the user can optionally implement two additional functions : **init()** and **close()** called before and after each map or reduce task. The user can also specify a "partition function" to send each key  $k_2$  generated in map phase to a specific reducer destination. The reducer destination may be computed using only a part of the input key  $k_2$ . The signature of the **partition** function is:

**partition:**  $k_2 \longrightarrow \text{Integer},$

<sup>1</sup>HDFS: Hadoop Distributed File System.

where the output of **partition** should be a positive number strictly smaller than the number of reducers. Hadoop's default **partition** function is based on "hashing" the whole input key  $k_2$ .

### 3 A MapReduce Skew Insensitive Join Algorithm

As stated in the introduction section, MapReduce hash based join algorithms presented in [4, 15] may be inefficient in the presence of highly skewed data[11] due to the fact that in Map function in these algorithms, all the key-value pairs  $(k_1, v_1)$  representing the same entry for the join attribute are sent to the same reducer (In Map phase, emitted key-value pairs  $(k_2, v_2)$ , key  $k_2$  is generated by only using join attribute values in the manner that all records with the same join attribute value will be forwarded to the same reducer).

To avoid the effect of repeated keys, Map user-defined function should generate distinct output keys  $k_2$  even for records having the same join attribute value. This is made possible by using a user defined partitioning function in Hadoop : the reducer destination for a key  $k_2$  can be computed from different parts of key  $k_2$  and not by a simple hashing of all input key  $k_2$ . To this end, we introduce, in this section, a join algorithm called MRFA-Join (MapReduce Frequency Adaptive Join) based on distributed histograms and a random redistribution of repeated join attribute values combined with an efficient technique of redistribution where only relevant data is redistributed across the network during the shuffle phase of reduce step. A cost analysis for MRFA-Join is also presented to give for each computation step, an upper bound of execution time in order to prove the strength of our approach.

In this section, we describe the implementation of MRFA-Join using Hadoop MapReduce framework as it is, without any modification. Therefore, the support for fault tolerance and load balancing in MapReduce and Distributed File System are preserved if possible: the inherent load imbalance due to repeated values must be handled efficiently by the join algorithm and not by the MapReduce framework.

To compute the join,  $R \bowtie S$ , of two relations (or datasets)  $R$  and  $S$ , we assume that input relations  $R$  and  $S$  are divided into blocks (splits) of data. These splits are stored in Hadoop Distributed File System (HDFS). These splits are also replicated on several nodes for reliability issues. Throughout this paper, for a relation  $T \in \{R, S\}$ , we use the following notations:

- $|T|$ : number of pages (or blocks of data) forming  $T$ ,
- $||T||$ : number of tuples (or records) in relation  $T$ ,
- $\bar{T}$ : the restriction (a fragment) of relation  $T$  which contains tuples which appear in the join result.  $||\bar{T}||$  is, in general, very small compared to  $||T||$ ,
- $T_i^{map}$ : the split(s) of relation  $T$  affected to mapper  $i$ ,
- $T_i^{red}$ : the split(s) of relation  $T$  affected to reducer  $i$ ,
- $\bar{T}_i$ : the split(s) of relation  $\bar{T}$  affected to mapper  $i$ ,
- $||T_i||$ : number of tuples in split  $T_i$ ,
- $Hist^{map}(T_i^{map})$ : Mapper's local histogram of  $T_i^{map}$ , i.e. the list of pairs  $(v, n_v)$  where  $v$  is a join attribute value and  $n_v$  its corresponding frequency in relation  $T_i^{map}$  on mapper  $i$ ,
- $Hist_i^{red}(T)$ : the fragment of global histogram of relation  $T$  on reducer  $i$ ,
- $Hist_i^{red}(T)(v)$  is the global frequency  $n_v$  of value  $v$  in relation  $T$ ,
- $HistIndex(R \bowtie S)$ : join attribute values that appear in both  $R$  and  $S$  and their corresponding three parameters: *Frequency\_index*, *Nb\_buckets1* and *Nb\_buckets2* used in communication templates,
- $c_{r/w}$ : read/write cost of a page of data from/to distributed file system (DFS),
- $c_{comm}$ : communication cost per page of data,

- $t_s^i$ : time to perform a simple search in a Hashtable on node  $i$ ,
- $t_h^i$ : time to add an entry to a Hashtable on node  $i$ ,
- $NB\_mappers$ : number of job mapper nodes,
- $NB\_reducers$ : number of job reducer nodes.

We will describe MRFA-Join algorithm while giving a cost analysis for each computation phase. Join computation in MRFA-Join proceeds in two MapReduce jobs:

- the first map-reduce job is performed to compute distributed histograms and to create randomized communication templates to redistribute only relevant data while avoiding the effect of data skew,
- the second one, is used to generate join output result by using communication templates carried out in the previous step.

In the following, we will describe MRFA-Join steps while giving an upper bound on the execution time of each MapReduce step. The  $O(\dots)$  notation only hides small constant factors: they only depend on program's implementation but neither on data nor on machine parameters. Data redistribution in MRFA-Join algorithm is the basis for efficient and scalable join processing while avoiding the effect of data skew in all the stages of join computation. MRFA-Join algorithm (see Algorithm 1) proceeds in 4 steps:

---

**Algorithm 1** MRFA-join algorithm workflow /\* See Appendix for detailed implementation \*/

---

- a.1**► Map phase: /\* To generate a tagged “Local histogram” for input relations \*/
- ▷ Each mapper  $i$  reads its assigned data splits (blocks) of relation  $R_i^{map}$  and  $S_i^{map}$  from the DFS
  - ▷ Extract the **join\_key** value from input relation's record.
  - ▷ Get a **tag** to identify source input relation.
  - ▷ Emit a couple  $((\text{join\_key}, \text{tag}), 1)$  /\* a **tagged join\_key** with a frequency 1 \*/
- Combine phase: To compute local frequencies for join\_key values in relations  $R_i^{map}$  and  $S_i^{map}$
- ▷ Each combiner, for each pair  $(\text{join\_key}, \text{tag})$  computes the sum of generated local frequencies associated to the join\_key value in each **tagged join\_key** generated in Map phase.
- Partition phase:
- ▷ for each emitted tagged join\_key, compute reducer destination according to only join\_key value.
- a.2**► Reduce phase: /\* To combine Shuffle's records and to create Global Join histogram index \*/
- ▷ Compute the global frequencies for only join\_key values present in both relations  $R$  and  $S$ .
  - ▷ Emit, for each join\_key, a couple  $(\text{join\_key}, (\text{frequency\_index}, \text{Nb\_buckets1}, \text{Nb\_buckets2}))$ .
- b.1**► Map phase:
- ▷ Each mapper reads join result global histogram index from DFS, and creates a local Hashtable.
  - ▷ Each mapper,  $i$ , reads its assigned data splits of input relations from DFS and generates randomized communication templates for records in  $R_i^{map}$  and  $S_i^{map}$  according to join\_key value and its corresponding frequency\_index in HashTable. In communication templates, only relevant records from  $R_i^{map}$  and  $S_i^{map}$  are emitted using hash or a randomized partition/replicate schema.
  - ▷ Emit relevant randomised tagged records from relations  $R_i^{map}$  and  $S_i^{map}$ .
- Partition phase:
- ▷ For each emitted tagged join\_key, compute reducer destination according to values of join\_key, and random reducer destination generated in Map phase;
- b.2**► Reduce phase: to combine Shuffle's output records and to generate join result.
- 

**a.1: Map phase to generate a tagged “local histogram” for input relations:**

In this step, each mapper  $i$  reads its assigned data splits (blocks) of relation  $R$  and  $S$  from distributed file system (DFS) and emits a couple  $(\langle K, \text{tag} \rangle, 1)$  for each record in  $R_i^{map}$  (resp.  $S_i^{map}$ ) where  $K$  is join key value and  $\text{tag}$  represents input relation tag. The cost of this step is :

$$Time(a.1.1) = O\left(\max_{i=1}^{NB\_mappers} c_{r/w} * (|R_i^{map}| + |S_i^{map}|) + \max_{i=1}^{NB\_mappers} (||R_i^{map}|| + ||S_i^{map}||)\right).$$

Emitted couples  $(\langle K, \text{tag} \rangle, 1)$  are then combined and partitioned using a user defined partitioning function by hashing only key part  $K$  and not the whole mapper tagged key  $\langle K, \text{tag} \rangle$ . The result of combine phase is then sent to reducers destination in the shuffle phase of the following reduce step. The cost of this step is at most:  $\text{Time}(a.1.2) =$

$$O\left(\max_{i=1}^{NB\_mappers} (||\text{Hist}^{map}(R_i^{map})|| * \log ||\text{Hist}^{map}(R_i^{map})|| + ||\text{Hist}^{map}(S_i^{map})|| * \log ||\text{Hist}^{map}(S_i^{map})||) + c_{comm} * (||\text{Hist}^{map}(R_i^{map})|| + ||\text{Hist}^{map}(S_i^{map})||)\right).$$

And the global cost of this step is:  $\text{Time}_{step_{a.1}} = \text{Time}(a.1.1) + \text{Time}(a.1.2)$ .

We recall that, in this step, only local histograms  $\text{Hist}^{map}(R_i^{map})$  and  $\text{Hist}^{map}(S_i^{map})$  are sorted and transmitted across the network and the sizes of these histograms are very small compared to the size of input relations  $R_i^{map}$  and  $S_i^{map}$  owing to the fact that, for a relation  $T$ ,  $\text{Hist}^{map}(T)$  contains only distinct entries of the form  $(v, n_v)$  where  $v$  is a join attribute value and  $n_v$  the corresponding frequency.

## a.2: Reduce phase to create join result global histogram index and randomized communication templates for relevant data:

At the end of shuffle phase, each reducer  $i$  will receive a fragment of  $\text{Hist}_i^{red}(R)$  (resp.  $\text{Hist}_i^{red}(S)$ ) obtained through hashing of distinct values of  $\text{Hist}^{map}(R_j^{map})$  (resp.  $\text{Hist}^{map}(S_j^{map})$ ) of each mapper  $j$ . Received  $\text{Hist}_i^{red}(R)$  and  $\text{Hist}_i^{red}(S)$  are then merged to compute global histogram  $\text{HistIndex}_i(R \bowtie S)$  on each reducer  $i$ .  $\text{HistIndex}(R \bowtie S)$  is used to compute randomized communication templates for only records associated to relevant join attribute values (i.e. values which will effectively be present in the join result).

In this step, each reducer  $i$ , computes the global frequencies for join attribute values which are present in both left and right relations and emits, for each join attribute  $K$ , an entry of the form :  $(K, \langle \text{Frequency\_index}(K), \text{Nb\_buckets1}(K), \text{Nb\_buckets2}(K) \rangle)$  where:

- $\text{Frequency\_index}(K) \in \{0, 1, 2\}$  will allow us to decide if, for a given relevant join attribute value  $K$ , the frequencies of tuples of relations  $R$  and  $S$  having the value  $K$  are greater (resp. smaller) than a defined threshold frequency  $f_0$ . It also permits us to choose dynamically the probe and the build relation for each value  $K$  of the join attribute. This choice reduces the global redistribution cost to a minimum.

For a given join attribute value  $K \in \text{HistIndex}_i(R \bowtie S)$ ,

$$\begin{cases} \text{Frequency\_index}(K)=0 & \text{If } \text{Hist}_i^{red}(R)(K) < f_0 \text{ and } \text{Hist}_i^{red}(S)(K) < f_0 \\ & \text{(i.e. values associated to low frequencies in both relations),} \\ \text{Frequency\_index}(K)=1 & \text{If } \text{Hist}_i^{red}(R)(K) \geq f_0 \text{ and } \text{Hist}_i^{red}(R)(K) \geq \text{Hist}_i^{red}(S)(K) \\ & \text{(i.e. Frequency in relation R is higher than those of S),} \\ \text{Frequency\_index}(K)=2 & \text{If } \text{Hist}_i^{red}(S)(K) \geq f_0 \text{ and } \text{Hist}_i^{red}(S)(K) > \text{Hist}_i^{red}(R)(K) \\ & \text{(i.e. Frequency in relation S is higher than those of R).} \end{cases}$$

- $\text{Nb\_buckets1}(K)$ : is the number of buckets used to partition records of relation associated to the highest frequency for join attribute value  $K$ ,
- $\text{Nb\_buckets2}(K)$ : is the number of buckets used to partition records of relation associated to the lowest frequency for join attribute value  $K$ .

For a join attribute value  $K$ , the number of buckets  $\text{Nb\_buckets1}(K)$  and  $\text{Nb\_buckets2}(K)$  are generated in a manner that each bucket will fit in reducer's memory. This makes the algorithm insensitive to the effect of data skew even for highly skewed input relations.

Figure 2 gives an example of communication templates used to partition data for  $\text{HistIndex}$  entry  $(K, \langle \text{Frequency\_index}(K), \text{Nb\_buckets1}(K), \text{Nb\_buckets2}(K) \rangle)$  corresponding to a join attribute

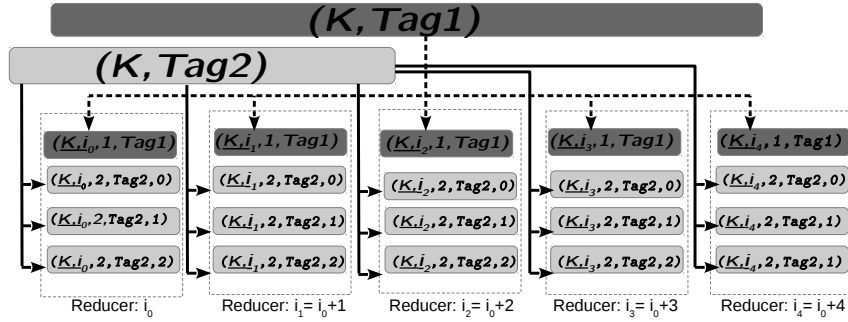


Figure 2: Generated buckets associated to a join key  $K$  corresponding to a high frequency where records from relation associated to  $Tag_1$  (i.e relation having the highest frequency) are partitioned into five buckets and those of relation associated to  $Tag_2$  are partitioned into three buckets.

$K$  associated to a high frequency, into small buckets. In this example, data associated to relation corresponding to  $Tag_1$  is partitioned into 5 buckets (i.e.  $Nb.buckets1(K) = 5$ ) where as those of relation corresponding to  $Tag_2$  is partitioned into 3 buckets (i.e.  $Nb.buckets2(K) = 3$ ). For these buckets, appropriate map keys are generated so that all records in each bucket of relation associated to  $Tag_1$  are forwarded to the same reducer holding all the buckets of relation associated to  $Tag_2$ . This partitioning guarantees that join tasks, are generated in a manner that the input data for each join task will fit in the memory of processing node and never exceed a user defined size, even for highly skewed data.

Using *HistIndex* information, each reducer  $i$ , has local knowledge of how relevant records of input relations will be redistributed in the next map phase. The global cost of this step is at most:  $Time_{step a.2} = O(\max_{i=1}^{NB-reducers} (||Hist_i^{red}(R)|| + ||Hist_i^{red}(S)||))$ . Note that,  $HistIndex(R \bowtie S) \equiv \cup_i (Hist_i^{red}(R) \cap Hist_i^{red}(S))$  and  $||HistIndex(R \bowtie S)||$  is very small compared to  $||Hist^{red}(R)||$  and  $||Hist^{red}(S)||$ .

To guarantee a perfect balancing of the load among processing nodes, communication templates are carried out jointly by all reducers (and not by a coordinator node) for only join attribute values which are present in join result : Each reducer deals with the redistribution of the data associated to a subset of relevant join attribute values.

### b.1: Map phase to create a local hash table and to redistribute relevant data using randomized communication templates:

In this step, each mapper  $i$  reads join result global histogram index, *HistIndex*, to create a local hash table in time:  $Time(b.1.1) = O(\max_{i=1}^{NB-mappers} t_h^i * ||HistIndex(R \bowtie S)||)$ .

Once local hash table is created on each mapper, input relations are then read from DFS, and each record is either discarded (if record's join key is not present in the local hash table) or routed to a designated random reducer destination using communication templates computed in step a.2 (Map phase details are described in Algorithm 6). The cost of this step is:

$$Time(b.1.2) = O\left(\max_{i=1}^{NB-mappers} (c_{r/w} * (|R_i^{map}| + |S_i^{map}|) + t_s^i * (||R_i^{map}|| + ||S_i^{map}||) + ||\bar{R}_i^{map}|| * \log ||\bar{R}_i^{map}|| + ||\bar{S}_i^{map}|| * \log ||\bar{S}_i^{map}|| + c_{comm} * (|\bar{R}_i^{map}| + |\bar{S}_i^{map}|))\right).$$

The term  $c_{r/w} * (|R_i^{map}| + |S_i^{map}|)$  is time to read input relations from DFS on each mapper

$i$ , the term  $t_s^i * (||R_i^{map}|| + ||S_i^{map}||)$  is the time to perform a hash table search for each input record,  $||\bar{R}_i^{map}|| * \log ||\bar{R}_i^{map}|| + ||\bar{S}_i^{map}|| * \log ||\bar{S}_i^{map}||$  is time to sort relevant data on mapper  $i$ , where as the term  $c_{comm} * (|\bar{R}_i^{map}| + |\bar{S}_i^{map}|)$  is time to communicate relevant data from mappers to reducers, using our communication templates described in step a.2. Hence the global cost of this step is:  $Time_{step_{b,1}} = Time(b.1.1) + Time(b.1.2)$ .

We recall that, in this step, only relevant data is emitted by mappers (which reduces communication cost in the shuffle step to a minimum) and records associated to high frequencies (those having a large effect on data skew) are redistributed according to an efficient dynamic partition/replicate schema to balance load among reducers and avoid the effect of data skew. However records associated to low frequencies (these records have no effect on data skew) are redistributed using hashing functions.

### b.2: Reduce phase to compute join result:

At the end of step b.1, each reducer  $i$  receives a fragment  $\bar{R}_i^{red}$  (resp.  $\bar{S}_i^{red}$ ) obtained through randomized hashing of  $\bar{R}_j^{map}$  (resp.  $\bar{S}_j^{map}$ ) of each mapper  $j$  and performs a local join of received data. This reduce phase is described in detail in Algorithm 8. The cost of this step is:

$$Time_{step_{b,2}} = O(\max_{i=1}^{NB\_reducers} (||\bar{R}_i^{red}|| + ||\bar{S}_i^{red}|| + c_{r/w} * |\bar{R}_i^{red} \bowtie \bar{S}_i^{red}|)).$$

The global cost of MRFA-Join is therefore the sum of the above four steps:

$$Time_{MRFA-Join} = Time_{step_{a,1}} + Time_{step_{a,2}} + Time_{step_{b,1}} + Time_{step_{b,2}}$$

Using hashing technique, the join computation of  $R \bowtie S$  requires at least the following lower bound :  $bound_{inf} =$

$$\Omega\left(\max_{i=1}^{NB\_mappers} ((c_{r/w} + c_{comm}) * (|R_i^{map}| + |S_i^{map}|) + ||R_i^{map}|| * \log ||R_i^{map}|| + ||S_i^{map}|| * \log ||S_i^{map}||) + \max_{i=1}^{NB\_reducers} (||R_i^{red}|| + ||S_i^{red}|| + c_{r/w} * |R_i^{red} \bowtie S_i^{red}|)\right),$$

where  $c_{r/w} * (|R_i^{map}| + |S_i^{map}|)$  is the cost of reading input relations from DFS on node  $i$ . The term  $||R_i^{map}|| * \log ||R_i^{map}|| + ||S_i^{map}|| * \log ||S_i^{map}||$  represents the cost to sort input relations records on map phase. The term  $c_{comm} * (|R_i^{map}| + |S_i^{map}|)$  represents the cost to communicate data from mappers to reducers, the term  $||R_i^{red}|| + ||S_i^{red}||$  is time to scan input relations on reducer  $i$  and  $c_{r/w} * |R_i^{red} \bowtie S_i^{red}|$  represents the cost to store reducer's  $i$  join result on the DFS.

MRFA-Join algorithm has asymptotic optimal complexity when:  $||HistIndex(R \bowtie S)||$

$$\leq \max\left(\max_{i=1}^{NB\_mappers} (||R_i^{map}|| * \log ||R_i^{map}||, |S_i^{map}| * \log ||S_i^{map}||), \max_{i=1}^{NB\_reducers} ||R_i^{red} \bowtie S_i^{red}||\right), \quad (1)$$

this is due to the fact that, all other terms in  $Time_{MRFA-Join}$  are bounded by those of  $bound_{inf}$ . Inequality 1 holds, in general, since  $HistIndex(R \bowtie S)$  contains only distinct values that appear in both relations  $R$  and  $S$ .

**Remark:** In practice, data imbalance related to the use of hashing functions can be due to:

- a bad choice of used hash function. This imbalance can be avoided by using the hashing techniques presented in the literature making it possible to distribute evenly the values of the join attribute with a very high probability [5],
- an intrinsic data imbalance which appears when some values of the join attribute appear more frequently than others. By definition a hash function maps tuples having the same join attribute values to the same processor. There is no way for a clever



hash function to avoid load imbalance that results from these repeated values [7]. But this case cannot arise here owing to the fact that histograms contain only distinct values of the join attribute and the hashing functions we use are always applied to histograms or applied to randomized keys.

## 4 Experiments

To evaluate the performance of MRFA-Join algorithm presented in this paper, we compared our algorithm to the best known solutions called respectively `ImprovedRepartitionJoin` and `StandardRepartitionJoin`. `ImprovedRepartitionJoin` was introduced by Blanas et al. in [4], where as `StandardRepartitionJoin` is the join algorithm provided in Hadoop framework's contributions. We ran a large series of experiments where 60 `Virtual Machines` (VMs) were randomly selected from our university cluster using OpenNubula software for VMs administration. Each Virtual Machine has the following characteristics: 1 Intel(R) Xeon@2.53GHz CPU, 4 Cores, 2GB of Memory and 100GB of Disk. Setting up a Hadoop cluster consisted of deploying each centralised entity (namenode and jobtracker) on a dedicated `Virtual Machine` and co-deploying datanodes and tasktrackers on the rest of VMs. The data replication parameter was fixed to three in the HDFS configuration file.

To study the effect of data skew on performance, join attribute values in the generated data have been chosen to follow a Zipf distribution [16] as it is the case in most database tests: Zipf factor was varied from 0 (for a uniform data distribution) to 1.0 (for a highly skewed data). Input relations size was fixed to 400M records for the right relation ( $\sim 40$ GB of data) and 10M of records for the left relation ( $\sim 1$ GB of data) and the join result varying from approximately 35M to 1700M records (corresponding respectively to about 7GB and 340GB of output data).

We noticed in all the tests and also those presented in Figure 3, that our MRFA-Join algorithm outperforms both `ImprovedRepartitionJoin` and `StandardRepartitionJoin` algorithms even for low or moderated skew. We recall that our algorithm requires the scan of input data twice. The first scan is performed for histogram processing and the second one for join processing. The cost analysis and tests performed showed that the overhead related to histogram processing is compensated by the gain in join processing since only relevant data (that appears in the join result) is emitted by mappers in the map phase which reduce considerably the amount of data transmitted over the network in shuffle phase (see Figure 4). Moreover, for skew factors varying from 0.6 to 1.0, both `ImprovedRepartitionJoin` and `StandardRepartitionJoin` jobs fail due to lack of memory. This is due to the fact that, in the reduce phase, all the records emitted by the mappers having the same join key are sent and processed by the same reducer which makes both `ImprovedRepartitionJoin` and `StandardRepartitionJoin` algorithms very sensitive to data skew and limits their scalability. This cannot occur in MRFA-Join owing to the fact that attribute values associated to high frequencies are forwarded to distinct reducers using randomised join attribute keys and not by a simple hashing of record's join key.

## 5 Conclusion and Future Work

In this paper, we have introduced the first skew-insensitive join algorithm, called *MRFA-Join*, using MapReduce, based on distributed histograms and randomised keys redistribution approach for highly skewed data. The detailed information provided by these histograms, allows us to reduce communication costs to only relevant data while guaranteeing perfect balancing processing due to the fact that all the generated join tasks and buffered data never exceed a user




Figure 3: Data skew effect on Hadoop join processing time




Figure 4: Data skew effect on the amount of data moved across the network during shuffle phase

defined size using threshold frequencies. This makes the algorithm scalable and outperforming existing MapReduce join algorithms which fail to handle skewed data whenever a join task cannot fit in the available node's memory. It is to be noted that *MRFA-Join* can also benefit from MapReduce underlying load balancing framework in a heterogeneous or a multi-user environment since *MRFA-Join* is implemented without any change in the MapReduce framework. Our experience with join operations shows that the overhead related to distributed histograms processing remains very small compared to the gain in performance and communication costs since only relevant data is processed or redistributed across the network.

We expect a higher gain related to histograms preprocessing in complex queries computation due to the fact that histograms can be used to reduce drastically the costs of communication and disk I/O of intermediate data by generating only relevant data for each sub-query. We will explore these aspects in the context of more complex and pipelined join queries.

## References

- [1] M. Bamha and G. Hains. Frequency-adaptive join for Shared Nothing machines. *Parallel and Distributed Computing Practices*, 2(3):333–345, 1999.
- [2] Mostafa Bamha. An optimal and skew-insensitive join and multi-join algorithm for distributed architectures. In *DEXA*, volume 3588 of *LNCs*, pages 616–625. Springer, 2005.

- [3] Mostafa Bamha and Gaétan Hains. A skew-insensitive algorithm for join and multi-join operation on Shared Nothing machines. In *DEXA*, volume 1873 of *LNCS*, pages 644–653. Springer, 2000.
- [4] Spyros Blanas, Jignesh M. Patel, Vuk Ercegovac, Jun Rao, Eugene J. Shekita, and Yuanyuan Tian. A comparison of join algorithms for log processing in MapReduce. In *SIGMOD*, pages 975–986. ACM, 2010.
- [5] J. Lawrence Carter and Mark N. Wegman. Universal Classes of Hash Functions. *Journal of Computer and System Sciences*, 18(2):143–154, 1979.
- [6] Jeffrey Dean and Sanjay Ghemawat. MapReduce: Simplified Data Processing on Large Clusters. In *OSDI*, pages 137–150. USENIX Association, 2004.
- [7] D. J. DeWitt, J. F. Naughton, D. A. Schneider, and S. Seshadri. Practical Skew Handling in Parallel Joins. In *VLDB*, pages 27–40, 1992.
- [8] Ralf Lämmel. Google’s MapReduce programming model – Revisited. *Science of Computer Programming*, 70(1):1–30, 2008.
- [9] Kyong-Ha Lee, Yoon-Joon Lee, Hyunsik Choi, Yon Dohn Chung, and Bongki Moon. Parallel Data Processing with MapReduce: A Survey. *ACM SIGMOD Record*, 40(4):11–20, 2011.
- [10] A. N. Mourad, R. J. T. Morris, A. Swami, and H. C. Young. Limits of parallelism in hash join algorithms. *Performance Evaluation*, 20(1/3):301–316, 1994.
- [11] Andrew Pavlo, Erik Paulson, Alexander Rasin, Daniel J. Abadi, David J. Dewitt, Samuel Madden, and Michael Stonebraker. A comparison of approaches to large-scale data analysis. In *SIGMOD*, pages 165–178. ACM, 2009.
- [12] D. Schneider and D. DeWitt. A Performance Evaluation of Four Parallel Join Algorithms in a Shared-Nothing Multiprocessor Environment. In *SIGMOD*. ACM, 1989.
- [13] M. Seetha and P. S. Yu. Effectiveness of parallel joins. *IEEE, Transactions on Knowledge and Data Enginneerings*, 2(4):410–424, 1990.
- [14] Tom White. *Hadoop – The Definitive Guide*. O’Reilly, second edition, 2010.
- [15] Hung-Chih Yang, Ali Dasdan, Ruey-Lung Hsiao, and D. Stott Parker. Map-reduce-merge: simplified relational data processing on large clusters. In *SIGMOD*, pages 1029–1040. ACM, 2007.
- [16] G. K. Zipf. *Human Behavior and the Principle of Least Effort: An Introduction to Human Ecology*. Adisson-Wesley, 1949.

## A Appendix: Implementation of MRFA-Join functions

---

**Algorithm 2** Map function /\* To generate local histograms values and tag input relation records \*/

---

```

map( $K$ : null,  $V$  : a record from a split of either relation  $R$  or  $S$ ) {
  ▷  $\text{relation\_tag} \leftarrow$  get relation tag from current relation split;
  ▷  $\text{join\_key} \leftarrow$  extract the join column from record  $V$  of relation  $R$ ;
  ▷ Emit ( $(\text{join\_key}, \text{relation\_tag})$ , 1);
}

```

---



---

**Algorithm 3** Combine function: /\* To compute local histogram’s frequencies for join\_key \*/

---

```

combine(Key  $K$ , List List_ $V$  ) {      /* List_ $V$  is the list of values “1” corresponding to the unique
                                     frequencies in relation  $R_i$  or  $S_i$  emitted by Mappers */
  ▷  $\text{frequency} \leftarrow$  sum of frequencies in List_ $V$ ;
  ▷ Emit ( $K, \text{frequency}$ );
}

```

---

---

**Algorithm 4** Partitioning function /\* Returns for, each composite key  $K=(\text{join\_key}, \text{relation\_tag})$  emitted in Map phase, an integer corresponding to destination reducer for the input key  $K$ . \*/

---

```

int partition(K: input key ){
  ▷ join_key ← K.join_key;          /* extracts join key part from input key  $K$  */
  ▷ Return (HashCode(join_key) % NB_reducers);
}

```

---



---

**Algorithm 5** Reduce function /\* To compute  $HistIndex(R \bowtie S)$  Global histogram index \*/

---

```

void reduce_init() {
  hash_index ← 0;          /* a flag to identify low frequencies records to redistribute using hashing */
  partition_index ← 1;     /* a flag to identify relation's records to partition */
  replicate_index ← 2;     /* a flag to identify relation's records to replicate */
  last_inner_key ← "";     /* to store the last processed key in inner relation */
  last_inner_frequency ← 0; /* to store the frequency of the last processed key in inner relation */
  /* THRESHOLD_FREQ: a user defined threshold frequency used for communication templates */
}

reduce(Key  $K$ , List List_V) { /* List_V: list of local frequencies of join_key in either  $R_i^{map}$  or  $S_i^{map}$  */
  ▷ join_key ← K.join_key;          /* extracts join key part from input key  $K$  */
  ▷ relation_tag ← K.relation_tag; /* extracts relation tag part from input key  $K$  */
  If (relation_tag corresponds to inner relation ) Then
    ▷ last_inner_key ← join_key;
    ▷ last_inner_frequency ← sum of frequencies in List_V;
  Else If (join_key = last_inner_key) Then
    ▷ frequency ← sum of frequencies in List_V ;
    If ((last_inner_frequency < THRESHOLD_FREQ) and (frequency < THRESHOLD_FREQ) Then
      ▷ Emit (join_key, (hash_index, 1, 1));
    ElseIf (last_inner_frequency ≥ frequency)
      ▷ Nb_buckets1 ← ⌈last_inner_frequency / THRESHOLD_FREQ⌉ ;
      ▷ Nb_buckets2 ← ⌈frequency / THRESHOLD_FREQ⌉;
      ▷ Emit (join_key, (partition_index, Nb_buckets1, Nb_buckets2));
    Else
      ▷ Nb_buckets1 ← ⌈frequency / THRESHOLD_FREQ⌉;
      ▷ Nb_buckets2 ← ⌈last_inner_frequency / THRESHOLD_FREQ⌉;
      ▷ Emit (join_key, (replicate_index, Nb_buckets1, Nb_buckets2));
    End If;
  End If;
End If;
}

```

---

---

**Algorithm 6** Map function: /\* To generate relevant randomized tagged records for input relations using *HistIndex* communication templates.\*/

---

```

void map_init() {
    inner_tag ← 1 ;           /* a tag to identify relation R records */
    outer_tag ← 2 ;          /* a tag to identify relation S records */
    hash_index ← 0 ;         /* a flag to identify hash based records */
    partition_index ← 1 ;    /* a flag to identify records to partition */
    replicate_index ← 2 ;    /* a flag to identify records to replicate */
    Read HistIndex( $R \bowtie S$ ): histogram index from DFS;
    Create a HashTable using join_key value, frequency's index and Nb_buckets of HistIndex( $R \bowtie S$ );
}

map( $K$ : null,  $V$  : a record from a split of either relation  $R$  or  $S$ ) {
    ▷ relation_tag ← get relation tag from current relation split;
    ▷ join_key ← extract the join column from record  $V$  of current input relation;
    If (join_key ∈ HashTable) Then /* To redistribute only relevant records */
        ▷ frequency_index ← HashTable(join_key).frequency_index;
        ▷ Nb_buckets1 ← HashTable(join_key).Nb_buckets1;
        ▷ Nb_buckets2 ← HashTable(join_key).Nb_buckets2;
        ▷ random_integer ← Generate_Random_Integer(join_key);
        If (frequency_index = hash_index) Then
            ▷ Emit ((join_key,-1,relation_tag),  $V$ ); /* for records, with low frequencies, to be hashed */
        ElseIf (((frequency_index = partition_index) and (relation_tag = inner_tag))
            or ((frequency_index = replicate_index) and (relation_tag=outer_tag)))
            ▷ random_dest ← (random_integer+SRAND(Nb_buckets1)) % Nb_buckets1;
            /* A random integer between 0 and Nb_buckets1 */
            ▷ flag_index ← partition_index ;
            ▷ Emit ((join_key,random_dest,(flag_index,relation_tag)),  $V$ );
        Else
            For (int i=0; i<Nb_buckets1; i++) Do
                ▷ random_dest ← (random_integer+i) % Nb_buckets1;
                ▷ flag_index ← replication_index ;
                ▷ bucket_dest ← i % Nb_buckets2; /* A random integer between 0 and Nb_buckets2 */
                ▷ Emit ((join_key,random_dest,(flag_index,relation_tag,bucket_dest)),  $V$ );
            End For;
        End If;
    End If;
}

```

---

**Algorithm 7** Partitioning function /\* Returns for each composite input key  $K = (\text{join\_key}, \text{random\_integer}, \text{DataTags})$  emitted in Map phase, an integer corresponding to destination reducer for key  $K$ . \*/

---

```

int partition( $K$ : input key ) {
    join_key ←  $K$ .join_key; /* extracts join key part from input key  $K$  */
    relation_tag ←  $K$ .relation_tag; /* extracts relation tag part from input key  $K$  */
    reducer_dest ←  $K$ .random_dest; /* extracts reducer destination number from input key  $K$  */
    If (reducer_dest ≠ -1) Then
        Return (reducer_dest % NB_reducers);
    Else
        Return (HashCode(join_key) % NB_reducers);
    End If;
}

```

---

---

**Algorithm 8** Reduce function: /\* To generate join result. \*/

---

```

void reduce.init() {
  last_key ← "" ;                               /* to store the last processed key */
  inner_relation_tag ← 1 ;                       /* a tag to identify Inner relation records */
  outer_relation_tag ← 2 ;                       /* a tag to identify Outer relation records */
  Array_buffer ← NULL ;                          /* an array list used to buffer records from one relation */
}
reduce(Key  $K$ , List List_V ) { /* List List_V: the list of records from either relation  $R$  or  $S$  */
  ▷ join_key ←  $K$ .join_key;                       /* extracts the join key part from input key  $K$  */
  ▷ relation_tag ←  $K$ .relation_tag;                /* extracts relation tag part from input key  $K$  */
  ▷ flag_index ←  $K$ .flag_index;                    /* extracts flag index part from input key  $K$  */
  If ((join_key = last_key) and (relation_tag ≠ flag_index)) Then
    For each record ( $x \in \text{List\_V}$ ) Do
      For each record ( $y \in \text{Array\_buffer}$ ) Do
        If (relation_tag = outer_relation_tag) Then
          ▷ Emit (NULL,  $x \oplus y$ );
        Else
          ▷ Emit (NULL,  $y \oplus x$ );
        End If ;
      End For ;
    End For ;
  Else
    ▷ Array_buffer.Clear();
    For each record ( $x \in \text{List\_V}$ ) Do
      ▷ Array_buffer.Add( $x$ );
    End For ;
    ▷ last_key ←  $K$ .join_key;
  End if
}

```

---